

Lecture 3

From last class: 50-60% of file system implemented

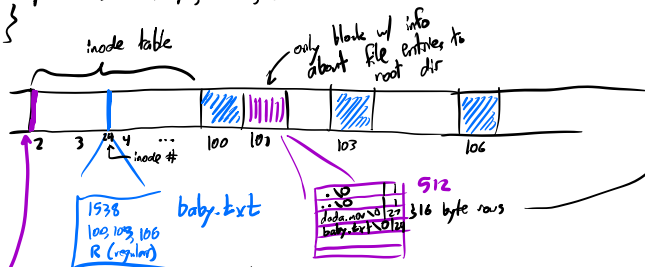
struct inode {

size_t size;

short blocknums [8];

char type; // reg. file or directory?

permissions, timestamps, owners etc.



struct dirent {
char name[14];
short inumber;
}

haven't done directories yet! → a special sort of file

for now: support root dir! → contents store filenames & associated inode nums

want to map baby.txt → 24 (the # inode it is)



inodes for usr & tmp → some payload struct!

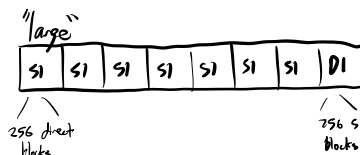
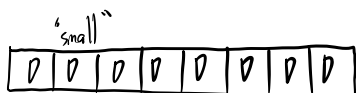
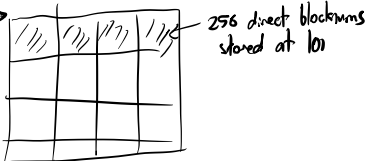


one problem not yet addressed: files only 4096 bytes!

assuming 8 blocknums are all direct — if file is big, ...



"indirect"



256 direct blocks

256 51 blocks

Major design principles

Modularity & Layering

→ special kind of modularity — org. of modules intermit in some hierarchical manner

symbolic link layer

→ each layer only opens the interface of layer above it

- absolute path name layer
- path name layer
- file name layer
- inode number layer
- file layer (how to aggregate sequences of blocks & store payloads)
- block layer

Naming & Name Resolution