

midway through dining philosophers problem

issue w/ current soln: deadlock

↳ introduced "permit" or "permissions slip"

assign 2 permission slips, or 4?

4: as high as possible w/o introducing deadlock

num forks = 1

not allowed to do the job of the thread manager & decide how threads are scheduled

unless you "know more" than the OS

int main(→) {

size_t permits = kNumForks - 1;

mutex forks[kNumForks], permitsLock;

thread philosophers[kNumPhilosophers];

grantPermission(size_t & permits, mutex & permitsLock) {

permitsLock.lock();

permits ++;

permitsLock.unlock();

}

waitForPermission()

// keep on looping until permit is available

// "willing to keep on polling permits until I'm sure its not zero

while (true) {

permitsLock.lock();

if (permits > 0) break;

permitsLock.unlock();

sleep_for(10); // 10 millisecond nap! allow other threads to use lock

}

permits --;

permitsLock.unlock();

we'd want something like sleep_until(permits > 0); ...

problem... this relies on "busy waiting"

also nothing special about 10 ms... we're making up numbers

want to be as sensible about CPU mgmt as possible

Condition Variables

condition_variable_any cv;

designed to monitor the truth/falsity of a condition you need in order to proceed

class condition_variable_any {

private:

// blissfully unaware

public:

void wait();

void notify_one(); // if there are multiple threads embedded in a wait call, pick 1 at random & let it proceed

void notify_all();

template < typename Pred > void wait(mutex & m, Pred pred);

now:

waitForPermission(→) {

m.lock();

alternatively: cv.wait(m, [& permits]() { return permits > 0; });

while (permits == 0) {

cv.wait(m); // don't want to go to sleep holding a lock! wait unlocks it for you

} // CPU intelligently renames lock when thread swapped off CPU - looks to the thread like they always have permission

permits++;

m.unlock();

}

grantPermission (→) {

m.lock();

permits++;

if (permits == 1) cv.notify_all();

m.unlock();

}

why not notify one? could have gone w/ just 2 permission slips...

think of them like push notifications to your phone

lock_guard<mutex> lg(m);

used in place of m.lock(); - useful for more complex threading examples

no way to leave the function w/o unlocking

don't have to explicitly unlock

Class Semaphore ← a class for the wait for / grant permission functionality above

public:

Semaphore(int count);

void wait(); ← decrements count

void signal(); ← increments & notifies all

private:

int count;

mutex m;

condition_variable any cv;

number of available resources

↳ a wait & notify mechanism

↳ a generic "resource counter" class

* do minimal amount of work needed to remove deadlock

* try to lock down resources for as narrow a time window as possible