

critical region: region of code we only want 1 thread accessing at a time

mutexes

mutual exclusion

"synchronization primitive"

like a lock on an open door

mutex m;

m.lock();

} critical
region

m.unlock(); // any other thread waiting on m.lock() can proceed at this pt. "taking the lock"

↑
does not depend on which thread arrived at lock first: a set, not a queue of waiting threads
typically named mutex var being Protected Lock;

pass 'm' mutex by reference!

ref(mlock) and mutex& mlock


↑
in thread() constructor

deadlock: all threads are 'stuck' (maybe at an m.lock()) but maybe for other reasons)

Dining Philosophers Problem

philosophers as array of threads

forks as array of mutexes



separate locks for all of your shared variables

*one way to check for race conditions in code: if you put in sleep calls & problems start to arise - they aren't creating problems, they're exposing them *

avoiding deadlock...

what about a stack of permits? only 2 philosophers can be eating at any given time

keep 4 permits - 1 fewer than would cause deadlock (e.g. w/ all 5 philosophers holding left forks)

mutexes: to prevent a thread from getting a stale old value of a variable

not necessary if all threads are only reading the variable!

necessary if writing is happening
my

goal: make critical region as small as possible while solution still works